

# COMP-598: Applied Machine Learning

## Mini-project #2: Text Classification

CMT submission for the report closing October 21, 11:59pm.

Kaggle submission closing October 21, 7:59pm (=11:59 UTC)

### Background:

For this project, you will participate in an in-class Kaggle competition on Text Classification. The goal is to devise a machine learning algorithm to analyze excerpts of conversations transcribed from interviews on the National Public Radio (NPR), and automatically classify them according to the topic (*AuthorInterviews*, *MovieInterviews*, *MusicInterviews* or *Interviews*). You will be able to download a training set, including labels, as well as a test set that does not include labels.

The competition, including the data, is available here (you must create a Kaggle account using your *mail.mcgill.ca* account.): <https://inclass.kaggle.com/c/dialogue-classification>

A description of the dataset is included: <https://inclass.kaggle.com/c/dialogue-classification/data>

Performance on the Kaggle Leaderboard will be calculated based on % of instances in the test set that are correctly classified.

The project should be completed in a group of 3. Remember: you must work with different team members on each mini-project. You can use the class discussion board on myCourses to find team members. A facebook group was also created to facilitate finding team members (see discussion board for the link).

### Requirements:

To participate in the competition, you must submit a list of predicted outputs for the test instances on the Kaggle website (see example file on Kaggle website for the file format). You can submit multiple prediction entries throughout the competition, and track your performance on the Kaggle Leaderboard. The test set is divided into two parts; one set (the public set) is used to update the scoreboard, one set (the private set) is used to calculate the final score of each participant.

To solve the problem, you should try the following methods (the 3<sup>rd</sup> one is optional):

- 1) A baseline algorithm consisting of Naïve Bayes (from lecture 5), fully implemented by your team.
- 2) A standard algorithm (from lectures 8-12), such as decision trees, nearest neighbor, or SVM, fully implemented by your team.
- 3) Any other machine learning method of your choice. Existing packages can be used, e.g. *scikit-learn*, if used appropriately, and with appropriate references in your report.

Your report should include results from all methods considered (so at least 1 *baseline* algorithm and 1 *standard* algorithm). For the Kaggle competition, you can submit results from your best performing method, from any of these categories.

**You are strongly encouraged to submit your report and code jointly as an IPython**

**Notebook.** Alternately, you can submit your report and code as two separate files, in which case the main text of the report should not exceed 5 pages. Figures, appendices and references can be in excess of this. The format should be double-column, 10pt font, min. 1" margins. You can use the standard IEEE conference format, e.g. [ewh.ieee.org/soc/dei/ceidp/docs/CEIDPFormat.doc](http://ewh.ieee.org/soc/dei/ceidp/docs/CEIDPFormat.doc).

Your report should include:

- Your team name (same as on Kaggle) in the title.
- A list of team members (enter them as "authors" on the submission website).
- Main text with the following sections:
  - o Introduction (overview of approach)
  - o Related work (previous literature related to this problem)
  - o Data pre-processing methods
  - o Feature design/selection methods
  - o Algorithm selection (for each of the categories *baseline/standard/advanced*)
  - o Optimization (if required for the algorithm)
  - o Parameter selection method (model order, learning rate, etc.)
  - o Testing and validation (detailed analysis of your results, outside of Kaggle)
  - o Discussion (pros/cons of your approach & methodology).
- When appropriate, use figures, tables and graphs to illustrate your work. Always include captions, axes labels, etc.
- Use appropriate referencing style throughout your report (with references listed in a separate section near the end, usually after the appendix).
- Before the references, add the following statement: "We hereby state that all the work presented in this report is that of the authors." Make sure this statement is truthful!
- Before the references, also add a section with a Statement of Contributions. Briefly describe the contributions of each team member towards each of the components of the project (e.g. defining the problem, developing the methodology, coding the solution, performing the data analysis, writing the report, etc.)
- Spell-check and proof-read carefully.

**Evaluation criteria:**

Marks will be attributed based on: 33% for performance on the private test set in the competition: 67% for the report including a clear description of the methods. The code will not be marked, but may be used to validate the other components.

For the competition, the performance grade will be calculated as follows: The top team, according to the score on the private test set, will receive 100%. A Random predictor, entered by the instructor, will score 0%. All other grades will be calculated according to interpolation of the Leaderboard scores between those two extremes.

For the report, the evaluation criteria include:

- Quality of review of related work

- Technical soundness of proposed methodology (feature selection, algorithms, optimization, validation plan)
- Clarity of methodology description, plots and figures.
- Overall organization and writing.

The same evaluation criteria will be used for peer-reviews and evaluation by TAs and instructor. The final report grade will be attributed 90% based on the assessment of TAs and instructor. The additional 10% is attributed following your participation in the peer-review process (i.e. for assessing reports of other groups).

Final grades and any late penalties will be attributed per team (i.e. all team members will get the same grade.)

You can discuss methods and technical issues with members of others teams, but you cannot share any code or data with other teams. Any team found to cheat (e.g. use external information, use resources without proper references) on either the code or report will receive a score of 0 for both the report and competition.

### **Submission instructions:**

Predictions on the test set must be submitted on Kaggle:

<https://inclass.kaggle.com/c/dialogue-classification/submit>

For the report, we will be using the same online conference management system:

[https://cmt.research.microsoft.com/COMP598\\_2015](https://cmt.research.microsoft.com/COMP598_2015)

You should use the same account as for the previous project, but submit to a new track, called “Mini-project #2”. The new report should be submitted as a “New Submission” (one per group), linking other team members as co-authors.

If you submit your code as a Notebook, you can submit a single file incorporating text and code. Otherwise, submit the report as a pdf, and the code as a “Supplementary” file.